

Bachelor of Science in Data Science



Course Code: DS3003

Course Name: Data Warehousing & Business Intelligence

Semester-Fall, 2024



Practical Project

**Building and Analysing a Near-Real-Time Data Warehouse
Prototype for METRO Shopping Store in Pakistan**

Total Marks: 100

Weight in grade: 15%

Due date: November 26, 2024, 11pm PST

1. Assessment task

The student has to design, implement, and analyse a near-real-time Data Warehouse (DW) prototype for METRO shopping store in Pakistan.

2. Project overview

METRO is one of the biggest superstores chains in Pakistan. The stores have thousands of customers and therefore it is important for the store to online analyse the shopping behaviour of their customers. Based on that the store can optimise their selling strategies e.g., giving promotions on different products.




Figure 1: An overview of METRO DW

Now, to make this analysis of shopping behaviour practical there is a need of building a near-real-time DW and customers' transactions from Data Sources (DSs) are required to reflect into DW as soon as they appear in DSs. The overview of METRO DW is presented in Figure 1. To build a near-real-time DW we need to implement a near-real-time ETL (Extraction, Transformation, and Loading) tools. Since the data generated by customers is incomplete as it required by DW, it needs to complete in the transformation layer of ETL. For example, enriching some information from Master Data (MD) as shown in Figure 2.




Figure 2: Enrichment example.

The crux of MESHJOIN is that with every loop step a new chunk of customers' transactions is read into main memory (**Hash table**) and MD partition in the disk-buffer is replaced by the new MD partition from the disk. Each of these chunks will remain in main memory for the time of one full MD cycle. The chunks therefore leave main memory in the order that they enter main memory and their time of residence in main memory is overlapping. This leads to the staggered processing pattern of MESHJOIN. In main memory, the incoming customers' data is organized in a queue, each chunk being one element of the queue. Figure 3 with four MD partitions shows a pictorial representation of the MESHJOIN operation: at each point in time, each chunk S_i in the queue has seen a larger number of partitions than the previous, and started at a later position in MD (except

for the case that the traversal of MD resets to the start of MD). The figure shows the moment when partition R_2 of MD is read into the disk-buffer but is not yet processed.

After loading the disk partition into the disk buffer, the algorithm probes each tuple of the disk buffer in the hash table. If a matching tuple is found, the algorithm generates the join output. After each iteration the algorithm removes the oldest chunk of customers' transactions from the hash table along with their pointers from the queue. This chunk is found at the end of the queue; its tuples were joined with the whole of MD and are thus completely processed now.

4. Star-schema

The star schema (which you will use in this project) is a data modelling technique that is used to map multidimensional decision support data into a relational database. Star-schema yields an easily implemented model for multidimensional data analysis while still preserving the relational structures on which the operational database is built.

The star schema represents aggregated data for specific business activities. Using the schema, one can create multiple aggregated data sources that will represent different aspects of business operations. For example, the aggregation may involve total sales by selected time periods, by products, by stores, and so on. Aggregated totals can be total product sold, total sales values by products, etc. The basic star schema has three main components: *facts*, *dimensions*, *attributes*, and *classification levels*. Figure 2 can help you to determine the right components for your star schema.

5. Implementation of Extended MESHJOIN

To implement MESHJOIN algorithm you will implement the following steps using Java Eclipse.

1. Read a segment of stream from TRANSACTIONS table as an input data into the hash table with their join attribute values in the queue.
2. Load next MD partitions from the both MD tables (customer and product) into the disk buffers. After the last partition the next partition to read will be the first partition in each table.
3. Perform join operation between the MD tuples and stream tuples.
4. If the join is successful, enrich the required information to the transaction tuple. It is important to note that the attribute TOTAL_SALE does not exist in both TRANSACTION and MD.
5. The transaction tuples will then be loaded into DW. Make sure the dimension tables will not have the duplication records.
6. Repeat steps 1 to 5 until you load all the data from TRANSACTIONS table to DW.

6. DW analysis

Once the entire data has been loaded into DW, you will be required to analyse your DW by applying following OLAP queries.

Q1. Top Revenue-Generating Products on Weekdays and Weekends with Monthly Drill-Down

Find the top 5 products that generated the highest revenue, separated by weekday and weekend sales, with results grouped by month for a specified year.

Q2. Trend Analysis of Store Revenue Growth Rate Quarterly for 2017

Calculate the revenue growth rate for each store on a quarterly basis for 2017.

Q3. Detailed Supplier Sales Contribution by Store and Product Category

For each store, show the total sales contribution of each supplier broken down by product category. The output should group results by store, then supplier, and then product category under each supplier.

Q4. Seasonal Analysis of Product Sales Using Dynamic Drill-Down

Present total sales for each product, drilled down by seasonal periods (Spring, Summer, Fall, Winter). This can help understand product performance across seasonal periods.

Q5. Store-Wise and Supplier-Wise Monthly Revenue Volatility

Calculate the month-to-month revenue volatility for each store and supplier pair. Volatility can be defined as the percentage change in revenue from one month to the next, helping identify stores or suppliers with highly fluctuating sales.

Q6. Top 5 Products Purchased Together Across Multiple Orders (Product Affinity Analysis)

Identify the top 5 products frequently bought together within a set of orders (i.e., multiple products purchased in the same transaction). This product affinity analysis could inform potential product bundling strategies.

Q7. Yearly Revenue Trends by Store, Supplier, and Product with ROLLUP

Use the ROLLUP operation to aggregate yearly revenue data by store, supplier, and product, enabling a comprehensive overview from individual product-level details up to total revenue per store. This query should provide an overview of cumulative and hierarchical sales figures.

Q8. Revenue and Volume-Based Sales Analysis for Each Product for H1 and H2

For each product, calculate the total revenue and quantity sold in the first and second halves of the year, along with yearly totals. This split-by-time-period analysis can reveal changes in product popularity or demand over the year.

Q9. Identify High Revenue Spikes in Product Sales and Highlight Outliers

Calculate daily average sales for each product and flag days where the sales exceed twice the daily average by product as potential outliers or spikes. Explain any identified anomalies in the report, as these may indicate unusual demand events.

Q10. Create a View REGION_STORE_QUARTERLY_SALES for Optimized Sales Analysis

Create a view named REGION_STORE_QUARTERLY_SALES that aggregates total quarterly sales by region and store, ordered by region and then by store name. This view allows quick retrieval of regional and store-specific trends across quarters, significantly improving query performance for regular sales analysis.

7. Tasks break-up

Following is list of tasks that you need to complete in this project.

1. Identifying appropriate dimension tables, fact table, and their attributes for the sales scenario presented in Figure 2. Based on that create a star-schema for DW with appropriate primary and foreign keys.

2. Implementing the MESHJOIN algorithm (using steps described in Section 6) in Java for successfully loading transactional data into DW after joining it with MD.
3. Applying of different analysis (described in Section 7) on DW using slicing, dicing, drill down, and materialising view concepts.
4. Writing a project report that should include project overview, schema for DW, MESHJOIN algorithm, any three shortcomings in Mesh Join, and what did you learn from the project?

8. What to submit

Each student has to submit the following files:

1. *Create-DW* –SQL script file to create star-schema for DW.
Note: your script should remove the pre-existed schema.
2. *Mesh-Join* – an Eclipse based Java project that implements MESHJOIN algorithm.
Note: You program should take the database credentials from the user at execution time.
3. *Project-Report* – a doc file containing all contents described in point 4 under the tasks break-up section.
4. *ReadMe* – a text file describing the step-by-step instructions to operate your project.

Note: all above files need to submit in a zipped folder named by your name, student ID e.g., Bilal_12i1234_Project.

9. When to submit?

Due date:

Late penalty: maximum late submissions time is 24 hours after the due date. In this case 15% late penalty will be applied.

10. Who to submit?

The project should be submitted through Google Class Room.

NOTE: Every student has to complete the project individually. Each student's project source and report materials should be unique and done by his/her own. All assessments will be assessed through Turnitin system and in case of finding any duplication or identical material **0 marks will be marked.**

Marking guide

Project Component	Marks
<i>Create-DW</i> – SQL script file to create star-schema for DW	/15
The script should create all dimensions and fact tables table in DW and if any table with same name exists already the script should drop that. It should also apply all primary and foreign keys on the right attributes.	
<i>Mesh-Join</i> – an Eclipse based Java project that implements MESHJOIN algorithm	/30
A Java file <i>mesh-Join</i> that should implement all three phases of ETL – it should extract records from TRANSACTIONS table, transform these with MD and then load these records to DW successfully.	

<i>Queries-DW</i> – SQL script file containing of all your OLAP queries	/30
The file should include OLAP queries for all tasks presented in Section 7.	
<i>Project-Report</i> – a doc file containing all contents described in point 4 under the task break-up section.	/20
Report must contain project overview, schema for DW, MESHJOIN algorithm, any three shortcomings in Mesh Join, and what did you learn from the project?	
<i>Read-Me</i> – a text file describing the step-by-step instructions to operate your project	/5
Read-Me file should contain a step-by-step guide to operate the project. In case of missing this file, 5 marks will be deducted.	
Late submission penalty	-/15%
TOTAL MARKS	/100

-----E N D-----